

Analytics Project

Your Name: IWAN Your G Number: G01236399

```
# Suppress dplyr summarise grouping warning messages
options(dplyr.summarise.inform = FALSE)

## Add R libraries here
library(tidyverse)
library(tidymodels)
library(vip)
library(rpart.plot)

# Load data
loans_df <- read_rds(url('https://gmubusinessanalytics.netlify.app/data/loan_data.rds'))
```

Data Analysis [50 Points]

In this section, you must think of at least 5 relevant questions that explore the relationship between `loan_default` and the other variables in the `loan_df` data set. The goal of your analysis should be discovering which variables drive the differences between customers who do and do not default on their loans.

You must answer each question and provide supporting data summaries with either a summary data frame (using `dplyr/tidyr`) or a plot (using `ggplot`) or both.

In total, you must have a minimum of 3 plots (created with `ggplot`) and 3 summary data frames (created with `dplyr`) for the exploratory data analysis section. Among the plots you produce, you must have at least 3 different types (ex. box plot, bar chart, histogram, scatter plot, etc...)

See the example question below.

Note: To add an R code chunk to any section of your project, you can use the keyboard shortcut **Ctrl + Alt + i** or the **insert** button at the top of your R project template notebook file.

Sample Question

Are there differences in loan default rates by loan purpose?

Answer: Yes, the data indicates that credit card and medical loans have significantly larger default rates than any other type of loan. In fact, both of these loan types have default rates at more than 50%. This is nearly two times the average default rate for all other loan types.

Summary Table

```
loans_df %>%
  group_by(loan_purpose) %>%
  summarise(n_customers = n(),
            customers_default = sum(loan_default == 'yes'),
            default_percent = 100 * mean(loan_default == 'yes'))
```

```
## # A tibble: 5 x 4
##   loan_purpose      n_customers customers_default default_percent
##   <fct>          <int>          <int>          <dbl>
## 1 debt_consolidation    1218            308            25.3
## 2 credit_card           879            470            53.5
## 3 medical               635            384            60.5
## 4 small_business        853            221            25.9
## 5 home_improvement      525            147            28.
```

Data Visualization

```
default_rates <- loans_df %>%
  group_by(loan_purpose) %>%
  summarise(n_customers = n(),
            customers_default = sum(loan_default == 'yes'),
            default_percent = 100 * mean(loan_default == 'yes'))

ggplot(data = default_rates,
       mapping = aes(x = loan_purpose,
                     y = default_percent)) +
  geom_bar(stat = 'identity',
          fill = '#006EA1',
          color = 'white') +
  labs(title = 'Loan Default Rate by Purpose of Loan',
       x = 'Loan Purpose',
       y = 'Default Percentage') +
  theme_light()
```



Question 1


Question: Are there differences in loan default rates by term?

Answer: Yes, there are difference in it. It precisely shown by the plot that the percentage of loan default in five years term has the higher rate compared to the three years term. However, the non default customers is likely to happen to happen in the three years term instead of in the five years term.

```
default_rates2_1 <- loans_df %>%  
  group_by(term, loan_default) %>%  
  summarise(n_customers = n()) %>%  
  mutate(percentage = (n_customers/sum(n_customers)) * 100)  
default_rates2_1
```

```
## # A tibble: 4 x 4  
## # Groups:   term [2]  
##   term      loan_default n_customers percentage  
##   <fct>      <fct>          <int>      <dbl>  
## 1 three_year yes             693        26.8  
## 2 three_year no            1895        73.2  
## 3 five_year  yes             837        55.0  
## 4 five_year  no             685        45.0
```

```
ggplot(data = default_rates2_1,  
  mapping = aes(x = term,  
                y = n_customers,  
                fill = loan_default)) +  
  geom_bar(stat = 'identity') +  
  geom_text(aes(label = round(percentage,  
                             digits = 2)),  
            size = 3,  
            position = position_stack(vjust = 0.5)) +  
  labs(title = 'Loan Default Rate by Term',  
        x = 'Term',  
        y = 'Number of Customers') +  
  theme_light()
```



Question 2

Question: Are there differences in loan default rates by current_job_years?

Answer: Yes, there are differences. It can clearly be explained by the line plot that the customers, which have the current job years in the middle year from 4 to 7 years, are likely not to default in the loan payment. However, there is a significant number of loan default customer spike for the customers in 8 years of work duration. It automatically explains that the most number of customers who likely to default in loan payment are customers which have the 8 years duration in current job years.


```
default_rates3 <- loans_df %>%
  group_by(current_job_years = as.factor(loans_df$current_job_years)) %>%
  summarise(n_customers = n(),
            customers_default = sum(loan_default == 'yes')) %>%
  mutate(default_percent = (customers_default/n_customers) * 100)
default_rates3
```

```
## # A tibble: 11 x 4
##   current_job_years n_customers customers_default default_percent
##   <fct>             <int>         <int>         <dbl>
## 1 0                 316           120           38.0
## 2 1                 303           117           38.6
## 3 2                 469           184           39.2
## 4 3                 389           140           36.0
## 5 4                 287           97            33.8
## 6 5                 314           107           34.1
```

##	7	6	173	61	35.3
##	8	7	166	55	33.1
##	9	8	132	62	47.0
##	10	9	144	55	38.2
##	11	10	1417	532	37.5

```
ggplot(data = default_rates3,
       mapping = aes(x = current_job_years,
                     y = default_percent,
                     group = 1)) +
  geom_line(color = "#0072B2") +
  geom_point(color = "#0072B2") +
  labs(title = "Line Plot of Default Percentage by Current Job Years",
       x = "Current Job Years",
       y = "Default Percent") +
  theme_light()
```

Line Plot of Default Percentage by Current Job Years



Question 3

Question: Do the default customers have the higher median of Loan Amount?

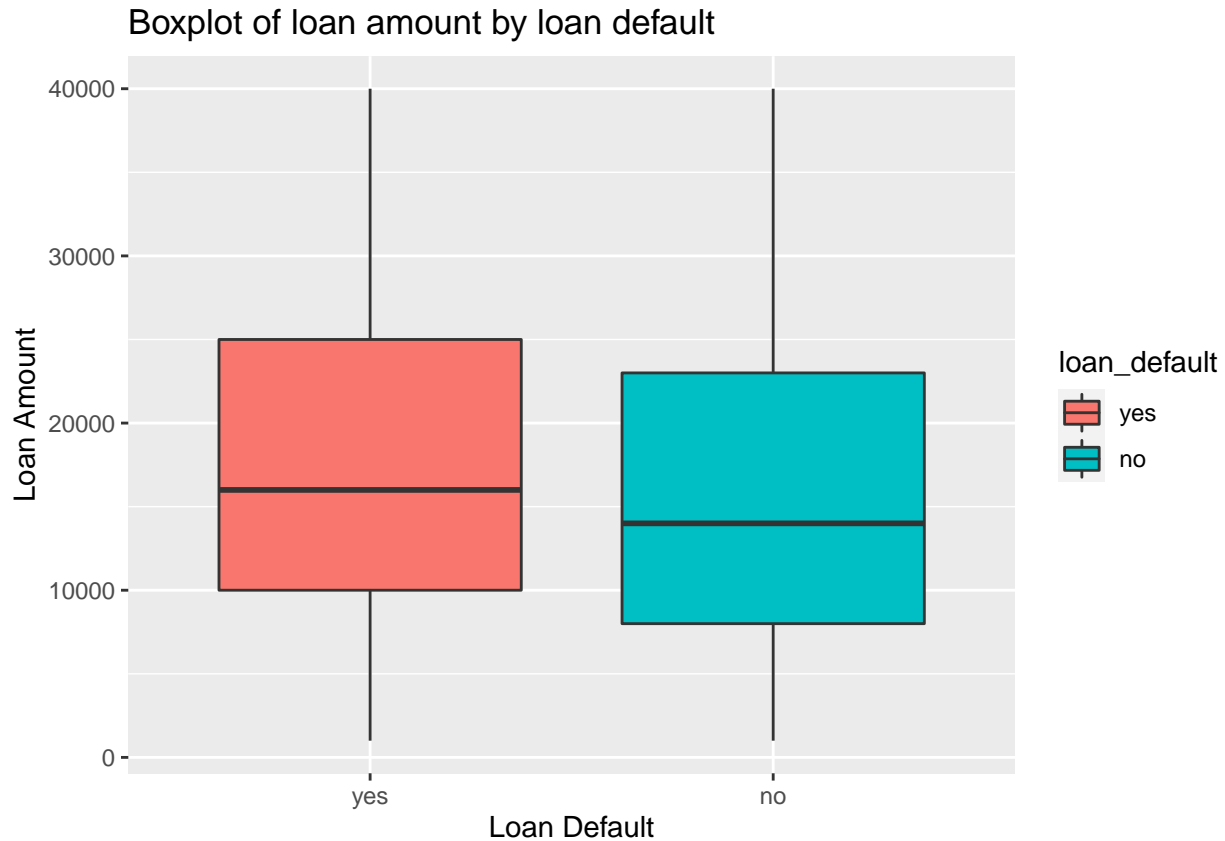
Answer: Yes, they do. The boxplot shows that the loan default has the higher median than the non default one. It can be concluded that the customers which likely to have the loan default status have the higher exposure of loan amount instead of the non default customer.

```
ggplot(data = loans_df,
       mapping = aes(x = loan_default,
```

```

    y = loan_amount,
    fill = loan_default)) +
geom_boxplot() +
labs(title = "Boxplot of loan amount by loan default",
     x = "Loan Default",
     y = "Loan Amount")

```



Question 4

Question: Are there differences in loan default based on the customers installment range?

Answer: Yes, there are. According to the installment range, it shows something interesting. The the loan default customers are likely to spike as the installment amount are getting higher. However, the non default percent of the customers are getting lower as the installment range are getting lower. Still, the average percentage of the loan default is lower than the non default customers.

```

installment_differences <- loans_df %>%
  mutate(installment_range = cut(installment,
                                breaks = c(0,
                                             275.00,
                                             422.00,
                                             490.00,
                                             664.00,
                                             1566.59),
                                labels = c("0 - 275.00",
                                           "275.01 - 422.00",

```

```

                                "422.01 - 490.00",
                                "490.01 - 664.00",
                                "664.01 - 1566.59")))) %>%

group_by(installment_range) %>%
summarise(n_customer = n(),
          customer_default = sum(loan_default == 'yes'),
          customer_non_default = sum(loan_default == 'no')) %>%
mutate(default_percent = (customer_default/n_customer) * 100,
       non_default_percent = (customer_non_default/n_customer) * 100)
installment_differences

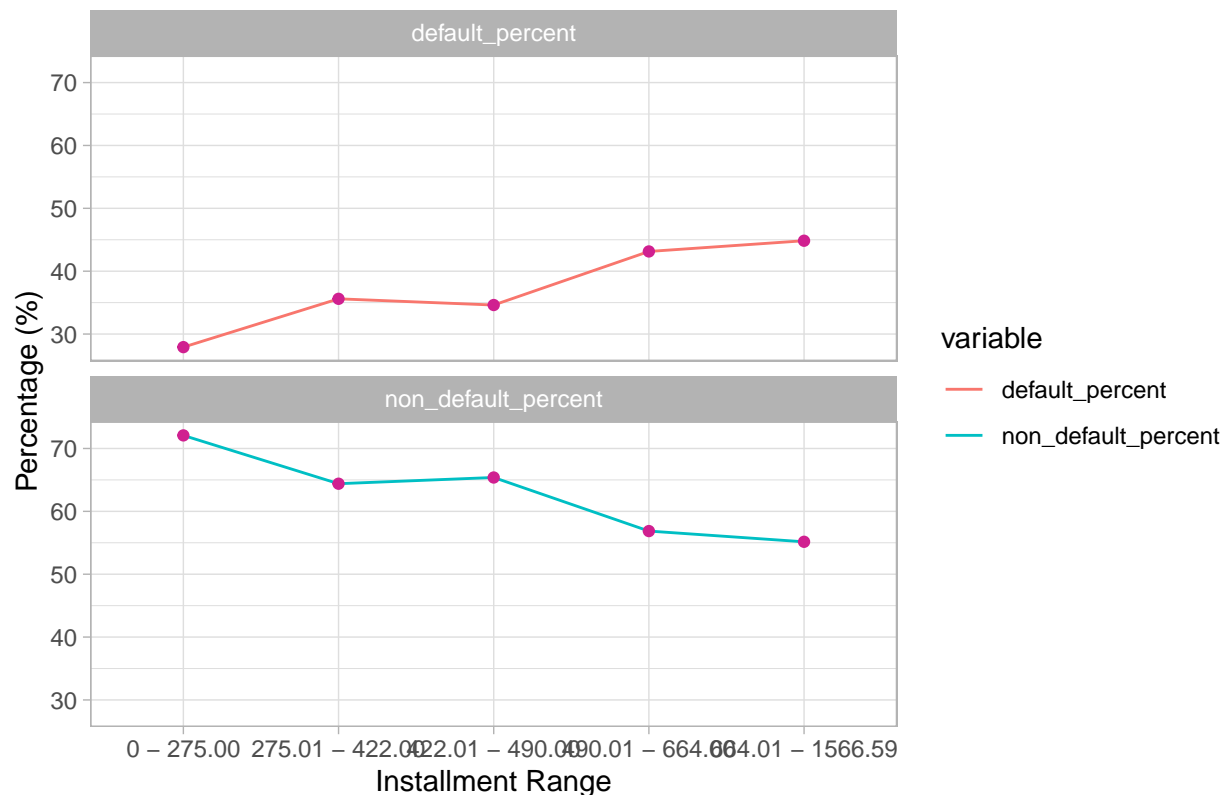
## # A tibble: 5 x 6
##   installment_ran~ n_customer customer_default customer_non_de~ default_percent
##   <fct>           <int>         <int>         <int>         <dbl>
## 1 0 - 275.00      1028          287          741          27.9
## 2 275.01 - 422.00 1028          366          662          35.6
## 3 422.01 - 490.00  312          108          204          34.6
## 4 490.01 - 664.00  714          308          406          43.1
## 5 664.01 - 1566.59 1028          461          567          44.8
## # ... with 1 more variable: non_default_percent <dbl>

#preparation for the plot
df_installment_differences <- installment_differences %>%
  select(installment_range,
         default_percent,
         non_default_percent) %>%
  gather(key = "variable",
         value = "value",
         -installment_range)

#plot the loan default based on the range installment
ggplot(df_installment_differences,
       aes(x = installment_range,
           y = value,
           group = variable,
           color = variable)) +
  geom_line() +
  geom_point(color="violetred") +
  facet_wrap(~ variable, ncol = 1) +
  labs(title = "Line Plot of Default Percentage by Installment Range",
       x = "Installment Range",
       y = "Percentage (%)") +
  theme_light()

```

Line Plot of Default Percentage by Installment Range



Question 5

Question: Are there differences in loan default rates by homeownership?

Answer: Yes, there are. According to summary, the customers, whose home ownership status is “rent”, have the higher percentage of the “mortgage” and the “own” status. Therefore, it can be concluded that the customers, who are renting house to live, have tendency to have the loan default status.

```
loans_df %>%
  group_by(homeownership) %>%
  summarise(n_customers = n(),
            customers_default = sum(loan_default == 'yes')) %>%
  mutate(default_percent = (customers_default/n_customers) * 100)
```

```
## # A tibble: 3 x 4
##   homeownership n_customers customers_default default_percent
##   <fct>         <int>         <int>         <dbl>
## 1 mortgage      1937             628           32.4
## 2 rent          1666             713           42.8
## 3 own           507             189           37.3
```

Predictive Modeling [75 Points]

In this section of the project, you will fit **two classification algorithms** to predict the response variable, `loan_default`. You should use all of the other variables in the `loans_df` data as predictor variables for each model.

You must follow the machine learning steps below.

The data splitting and feature engineering steps should only be done once so that your models are using the same data and feature engineering steps for training.

- Split the `loans_df` data into a training and test set (remember to set your seed)
- Specify a feature engineering pipeline with the `recipes` package
 - You can include steps such as skewness transformation, dummy variable encoding or any other steps you find appropriate
- Specify a `parsnip` model object
 - You may choose from the following classification algorithms:
 - * Logistic Regression
 - * LDA
 - * QDA
 - * KNN
 - * Decision Tree
 - * Random Forest
- Package your recipe and model into a workflow
- Fit your workflow to the training data
 - If your model has hyperparameters:
 - * Split the training data into 5 folds for 5-fold cross validation using `vfold_cv` (remember to set your seed)
 - * Perform hyperparameter tuning with grid search using the `grid_regular()` function
 - * Hyperparameter tuning can take a significant amount of computing time. To minimize this, use a maximum of 3 levels within your `grid_regular()` function
 - * Select the best model with `select_best()` and finalize your workflow
- Evaluate model performance on the test set by plotting an ROC curve using `autoplot()` and calculating the area under the ROC curve on your test data

Split the ‘loans_df’ data into a training and test set (remember to set your seed)

```
#1. Split the `loans_df` data into a training and test set (remember to set your seed)
#-----
set.seed(4500)
loans_split <- initial_split(loans_df, prop = 0.75,
                             strata = loan_default)

loans_training <- loans_split %>% training()
loans_test <- loans_split %>% testing()
```

Specify a feature engineering pipeline

```
#2. Specify a feature engineering pipeline with the `recipes` package
#-----
loans_recipe <- recipe(loan_default ~ .,
                        data = loans_training) %>%
  step_YeoJohnson(all_numeric(),
                    -all_outcomes()) %>%
  step_normalize(all_numeric(),
                  -all_outcomes()) %>%
  step_dummy(all_nominal(),
```

```

      -all_outcomes())

loans_recipe %>%
  prep() %>%
  bake(new_data = loans_training)

## # A tibble: 3,083 x 20
##   loan_amount installment interest_rate annual_income current_job_yea~
##   <dbl>         <dbl>         <dbl>         <dbl>         <dbl>
## 1      1.56         1.38         1.40         0.941         -1.00
## 2     -0.550       -0.741        0.0985       -0.293         1.10
## 3     -0.216       -0.167        1.95         0.192         1.10
## 4      1.56         1.38         1.60         1.99         -1.74
## 5      0.958        0.602        0.215        1.33        -0.126
## 6     -1.23       -1.22       -0.860        0.108       -0.401
## 7      0.0173      0.0659      0.0280       -0.860         1.10
## 8      1.62         1.31         0.752       -1.75         1.10
## 9     -0.899       -0.675        1.80        -1.32       -0.126
## 10     1.62         0.993       -1.09         2.24         1.10
## # ... with 3,073 more rows, and 15 more variables: debt_to_income <dbl>,
## #   total_credit_lines <dbl>, years_credit_history <dbl>, loan_default <fct>,
## #   loan_purpose_credit_card <dbl>, loan_purpose_medical <dbl>,
## #   loan_purpose_small_business <dbl>, loan_purpose_home_improvement <dbl>,
## #   application_type_joint <dbl>, term_five_year <dbl>,
## #   homeownership_rent <dbl>, homeownership_own <dbl>,
## #   missed_payment_2_yr_no <dbl>, history_bankruptcy_no <dbl>,
## #   history_tax_liens_no <dbl>

```

Model 1

```

#3. Specify a `parsnip` model object
#-----
# Model Specification
#-----
# Logistic Regression
#-----
logistic_model <- logistic_reg() %>%
  set_engine('glm') %>%
  set_mode('classification')
logistic_model

## Logistic Regression Model Specification (classification)
##
## Computational engine: glm

#4. Package your recipe and model into a workflow
#-----
# Create a Workflow for logistic regression
#-----
loans_wf <- workflow() %>%
  add_model(logistic_model) %>%
  add_recipe(loans_recipe)
loans_wf

```

```

## == Workflow =====
## Preprocessor: Recipe
## Model: logistic_reg()
##
## -- Preprocessor -----
## 3 Recipe Steps
##
## * step_YeoJohnson()
## * step_normalize()
## * step_dummy()
##
## -- Model -----
## Logistic Regression Model Specification (classification)
##
## Computational engine: glm
#5. Fit your workflow to the training data
#-----
# Fit the Model to the logistic regression model
#-----
# Fit the workflow to the training data
loans_logistic_fit <- loans_wf %>%
  fit(data = loans_training)
loans_logistic_fit

## == Workflow [trained] =====
## Preprocessor: Recipe
## Model: logistic_reg()
##
## -- Preprocessor -----
## 3 Recipe Steps
##
## * step_YeoJohnson()
## * step_normalize()
## * step_dummy()
##
## -- Model -----
##
## Call: stats::glm(formula = .y ~ ., family = stats::binomial, data = data)
##
## Coefficients:
##              (Intercept)              loan_amount
##                9.0680                27.1336
##            installment            interest_rate
##           -24.4930            -2.7640
##        annual_income            current_job_years
##           -0.2113                0.1001
##        debt_to_income            total_credit_lines
##           -0.3291                0.2564
##    years_credit_history    loan_purpose_credit_card
##           -0.1450            -1.0496
##    loan_purpose_medical    loan_purpose_small_business
##           -1.6486                0.3135
## loan_purpose_home_improvement    application_type_joint
##                0.1656            -0.6748


```

```
##           term_five_year           homeownership_rent
##           -15.5329           -0.6933
##           homeownership_own           missed_payment_2_yr_no
##           -0.5437           0.1715
##           history_bankruptcy_no           history_tax_liens_no
##           -0.1943           -0.6532
##
## Degrees of Freedom: 3082 Total (i.e. Null); 3063 Residual
## Null Deviance: 4071
## Residual Deviance: 773.9 AIC: 813.9

# Exploring our Trained Model
# Extract the trained model from our workflow fit
loans_trained_model <- loans_logistic_fit %>%
  pull_workflow_fit()
loans_trained_model

## parsnip model object
##
## Fit time: 30ms
##
## Call: stats::glm(formula = .y ~ ., family = stats::binomial, data = data)
##
## Coefficients:
##           (Intercept)           loan_amount
##           9.0680           27.1336
##           installment           interest_rate
##           -24.4930           -2.7640
##           annual_income           current_job_years
##           -0.2113           0.1001
##           debt_to_income           total_credit_lines
##           -0.3291           0.2564
##           years_credit_history           loan_purpose_credit_card
##           -0.1450           -1.0496
##           loan_purpose_medical           loan_purpose_small_business
##           -1.6486           0.3135
## loan_purpose_home_improvement           application_type_joint
##           0.1656           -0.6748
##           term_five_year           homeownership_rent
##           -15.5329           -0.6933
##           homeownership_own           missed_payment_2_yr_no
##           -0.5437           0.1715
##           history_bankruptcy_no           history_tax_liens_no
##           -0.1943           -0.6532
##
## Degrees of Freedom: 3082 Total (i.e. Null); 3063 Residual
## Null Deviance: 4071
## Residual Deviance: 773.9 AIC: 813.9

# Variable Importance
vip(loans_trained_model)
```



Predicted Categories

```
predictions_categories <- predict(loans_logistic_fit,
                                  new_data = loans_test)
```

predictions_categories

```
## # A tibble: 1,027 x 1
##   .pred_class
##   <fct>
## 1 no
## 2 yes
## 3 no
## 4 no
## 5 no
## 6 yes
## 7 yes
## 8 yes
## 9 no
## 10 no
## # ... with 1,017 more rows
```

obtain the estimated probabilities for each category of our response variable

```
predictions_probabilities <- predict(loans_logistic_fit,
                                     new_data = loans_test,
                                     type = 'prob')
```

predictions_probabilities

```
## # A tibble: 1,027 x 2
```

```
##      .pred_yes .pred_no
##      <dbl>    <dbl>
## 1  0.00131    0.999
## 2  0.552      0.448
## 3  0.00318    0.997
## 4  0.000143   1.00
## 5  0.0738     0.926
## 6  0.992      0.00753
## 7  0.996      0.00405
## 8  0.989      0.0114
## 9  0.00570    0.994
## 10 0.00101    0.999
## # ... with 1,017 more rows
```

combine the results from the Predicted Categories and the estimated probabilities with the true response

```
test_results <- loans_test %>%
  select(loan_default) %>%
  bind_cols(predictions_categories) %>%
  bind_cols(predictions_probabilities)
test_results
```

```
## # A tibble: 1,027 x 4
##   loan_default .pred_class .pred_yes .pred_no
##   <fct>        <fct>        <dbl>    <dbl>
## 1 no          no          0.00131  0.999
## 2 yes         yes         0.552    0.448
## 3 no          no          0.00318  0.997
## 4 no          no          0.000143 1.00
## 5 no          no          0.0738   0.926
## 6 yes         yes         0.992    0.00753
## 7 yes         yes         0.996    0.00405
## 8 yes         yes         0.989    0.0114
## 9 no          no          0.00570  0.994
## 10 no         no          0.00101  0.999
## # ... with 1,017 more rows
```

#6. Evaluate model performance on the test set by plotting an ROC curve using `autoplot()` and calculating the area under the curve

#-----

Evaluate the logistic regression model

#-----

Exploring Performance Metrics

Confusion Matrix

```
conf_mat(test_results,
  truth = loan_default,
  estimate = .pred_class)
```

```
##           Truth
## Prediction yes  no
##           yes 354 16
##           no  28 629
```

F1 Score


```
f_meas(test_results,
  truth = loan_default,
  estimate = .pred_class)
```

```
## # A tibble: 1 x 3
##   .metric .estimator .estimate
##   <chr>   <chr>       <dbl>
## 1 f_meas binary       0.941

# ROC Curve
roc_curve(test_results,
          truth = loan_default,
          estimate = .pred_yes)

## # A tibble: 1,029 x 3
##   .threshold specificity sensitivity
##   <dbl>         <dbl>         <dbl>
## 1 -Inf           0             1
## 2  1.16e-9        0             1
## 3  2.08e-9        0.00155       1
## 4  2.67e-9        0.00310       1
## 5  1.64e-8        0.00465       1
## 6  4.58e-8        0.00620       1
## 7  6.43e-8        0.00775       1
## 8  1.18e-7        0.00930       1
## 9  1.70e-7        0.0109        1
## 10 1.95e-7         0.0124        1
## # ... with 1,019 more rows

# plot the ROC Curve
roc_curve(test_results,
          truth = loan_default,
          estimate = .pred_yes) %>%
  autoplot()
```



```
# Area Under the ROC Curve
```

```
roc_auc(test_results,
        truth = loan_default, .pred_yes)
```

```
## # A tibble: 1 x 3
##   .metric .estimator .estimate
##   <chr>   <chr>      <dbl>
## 1 roc_auc binary      0.989
```

```
# Creating Custom Metric Sets
```

```
# calculate the accuracy and F1 from my results data frame
```

```
my_metrics <- metric_set(accuracy,
                        f_meas)
```

```
my_metrics(test_results,
          truth = loan_default,
          estimate = .pred_class)
```

```
## # A tibble: 2 x 3
##   .metric .estimator .estimate
##   <chr>   <chr>      <dbl>
## 1 accuracy binary      0.957
## 2 f_meas  binary      0.941
```

```
# Automating the Process
```

```
last_fit_model <- loans_wf %>%
  last_fit(split = loans_split)
last_fit_model
```



```
## # Resampling results
## # Monte Carlo cross-validation (0.75/0.25) with 1 resamples
## # A tibble: 1 x 6
##   splits      id      .metrics    .notes    .predictions    .workflow
##   <list>      <chr>    <list>    <list>    <list>          <list>
## 1 <split [3.1K~ train/test ~ <tibble [2 ~ <tibble [0~ <tibble [1,027 ~ <workflo~


# obtain the metrics on the test set (accuracy and roc_auc by default)
last_fit_model %>%
  collect_metrics()

## # A tibble: 2 x 3
##   .metric .estimator .estimate
##   <chr>   <chr>      <dbl>
## 1 accuracy binary      0.957
## 2 roc_auc  binary      0.989

# obtain a data frame with test set results
last_fit_results <- last_fit_model %>%
  collect_predictions()
last_fit_results

## # A tibble: 1,027 x 6
##   id      .pred_yes .pred_no .row .pred_class loan_default
##   <chr>      <dbl>    <dbl> <int> <fct>        <fct>
## 1 train/test split 0.00131 0.999     3 no         no
## 2 train/test split 0.552    0.448     4 yes        yes
## 3 train/test split 0.00318 0.997     5 no         no
## 4 train/test split 0.000143 1.00    10 no         no
## 5 train/test split 0.0738   0.926    28 no         no
## 6 train/test split 0.992    0.00753 33 yes        yes
## 7 train/test split 0.996    0.00405 34 yes        yes
## 8 train/test split 0.989    0.0114 48 yes        yes
## 9 train/test split 0.00570 0.994    52 no         no
## 10 train/test split 0.00101 0.999    69 no         no
## # ... with 1,017 more rows

# make an ROC plot
last_fit_results %>%
  roc_curve(truth = loan_default,
            estimate = .pred_yes) %>%
  autoplot()
```



Model 2

```
#-----
#3. Specify a `parsnip` model object
#-----
# Model Specification
#-----
# Decision Trees
#-----
tree_model <- decision_tree(cost_complexity = tune(),
                           tree_depth = tune(),
                           min_n = tune()) %>%
  set_engine('rpart') %>%
  set_mode('classification')
tree_model

## Decision Tree Model Specification (classification)
##
## Main Arguments:
##   cost_complexity = tune()
##   tree_depth = tune()
##   min_n = tune()
##
## Computational engine: rpart
```

```
#4. Package your recipe and model into a workflow
```

```
#-----
```

```
# Create a Workflow for decision tree
```

```
#-----
```

```
tree_workflow <- workflow() %>%  
  add_model(tree_model) %>%  
  add_recipe(loans_recipe)  
tree_workflow
```

```
## == Workflow =====
```

```
## Preprocessor: Recipe
```

```
## Model: decision_tree()
```

```
##
```

```
## -- Preprocessor -----
```

```
## 3 Recipe Steps
```

```
##
```

```
## * step_YeoJohnson()
```

```
## * step_normalize()
```

```
## * step_dummy()
```

```
##
```

```
## -- Model -----
```

```
## Decision Tree Model Specification (classification)
```

```
##
```

```
## Main Arguments:
```

```
##   cost_complexity = tune()
```

```
##   tree_depth = tune()
```

```
##   min_n = tune()
```

```
##
```

```
## Computational engine: rpart
```

```
# Create a grid of hyperparameter values to test
```

```
tree_grid <- grid_regular(cost_complexity(),  
                          tree_depth(),  
                          min_n(),  
                          levels = 3)
```

```
# View grid
```

```
tree_grid
```

```
## # A tibble: 27 x 3
```

```
##   cost_complexity tree_depth min_n  
##           <dbl>      <int> <int>
```

```
## 1  0.0000000001         1     2
```

```
## 2  0.00000316           1     2
```

```
## 3  0.1                 1     2
```

```
## 4  0.0000000001         8     2
```

```
## 5  0.00000316           8     2
```

```
## 6  0.1                 8     2
```

```
## 7  0.0000000001        15     2
```

```
## 8  0.00000316          15     2
```

```
## 9  0.1                15     2
```

```
## 10 0.0000000001         1    21
```

```
## # ... with 17 more rows
```

```

# Create folds for cross validation on the training data set
## These will be used to tune model hyperparameters
set.seed(4500)
loans_folds <- vfold_cv(loans_training, v = 5)
loans_folds

## # 5-fold cross-validation
## # A tibble: 5 x 2
##   splits          id
##   <list>         <chr>
## 1 <split [2.5K/617]> Fold1
## 2 <split [2.5K/617]> Fold2
## 3 <split [2.5K/617]> Fold3
## 4 <split [2.5K/616]> Fold4
## 5 <split [2.5K/616]> Fold5

## Tune decision tree workflow
set.seed(4500)
tree_tuning <- tree_workflow %>%
  tune_grid(resamples = loans_folds,
            grid = tree_grid)
tree_tuning

## # Tuning results
## # 5-fold cross-validation
## # A tibble: 5 x 4
##   splits          id   .metrics      .notes
##   <list>         <chr> <list>      <list>
## 1 <split [2.5K/617]> Fold1 <tibble [54 x 7]> <tibble [0 x 1]>
## 2 <split [2.5K/617]> Fold2 <tibble [54 x 7]> <tibble [0 x 1]>
## 3 <split [2.5K/617]> Fold3 <tibble [54 x 7]> <tibble [0 x 1]>
## 4 <split [2.5K/616]> Fold4 <tibble [54 x 7]> <tibble [0 x 1]>
## 5 <split [2.5K/616]> Fold5 <tibble [54 x 7]> <tibble [0 x 1]>

## Show the top 5 best models based on roc_auc metric
tree_tuning %>% show_best('roc_auc')## Select best model based on roc_auc

## # A tibble: 5 x 9
##   cost_complexity tree_depth min_n .metric .estimator  mean    n std_err
##   <dbl>          <int> <int> <chr>   <chr>      <dbl> <int>  <dbl>
## 1  0.0000000001      15    40 roc_auc binary    0.964     5 0.00299
## 2  0.00000316       15    40 roc_auc binary    0.964     5 0.00299
## 3  0.0000000001      15    21 roc_auc binary    0.964     5 0.00357
## 4  0.00000316       15    21 roc_auc binary    0.964     5 0.00357
## 5  0.0000000001      8    40 roc_auc binary    0.959     5 0.00305
## # ... with 1 more variable: .config <fct>

best_tree <- tree_tuning %>%
  select_best(metric = 'roc_auc')

# View the best tree parameters
best_tree

## # A tibble: 1 x 4
##   cost_complexity tree_depth min_n .config
##   <dbl>          <int> <int> <fct>

```

```
## 1      0.0000000001      15      40 Model25
```

```
# Finalize Workflow
```

```
final_tree_workflow <- tree_workflow %>%
```

```
  finalize_workflow(best_tree)
```

```
final_tree_workflow
```

```
## == Workflow =====
## Preprocessor: Recipe
## Model: decision_tree()
##
## -- Preprocessor -----
## 3 Recipe Steps
##
## * step_YeoJohnson()
## * step_normalize()
## * step_dummy()
##
## -- Model -----
## Decision Tree Model Specification (classification)
##
## Main Arguments:
##   cost_complexity = 1e-10
##   tree_depth = 15
##   min_n = 40
##
## Computational engine: rpart
```

```
#5. Fit your workflow to the training data
```

```
#-----
```

```
# Fit the Model to the decision tree model
```

```
#-----
```

```
# Fit the workflow to the training data
```

```
tree_wf_fit <- final_tree_workflow %>%
```

```
  fit(data = loans_training)
```

```
tree_wf_fit
```

```
## == Workflow [trained] =====
## Preprocessor: Recipe
## Model: decision_tree()
##
## -- Preprocessor -----
## 3 Recipe Steps
##
## * step_YeoJohnson()
## * step_normalize()
## * step_dummy()
##
## -- Model -----
## n= 3083
##
## node), split, n, loss, yval, (yprob)
##      * denotes terminal node
##
##      1) root 3083 1148 no (0.37236458 0.62763542)
##      2) interest_rate>=0.691997 702      0 yes (1.00000000 0.00000000) *
```

```
##      3) interest_rate< 0.691997 2381 446 no (0.18731625 0.81268375)
##      6) interest_rate>=-0.3606202 1287 446 no (0.34654235 0.65345765)
##     12) term_five_year>=0.5 473 234 yes (0.50528541 0.49471459)
##     24) installment>=-0.8083194 449 210 yes (0.53229399 0.46770601)
##     48) loan_amount< 1.59353 411 177 yes (0.56934307 0.43065693)
##     96) installment>=0.9637937 64 5 yes (0.92187500 0.07812500) *
##     97) installment< 0.9637937 347 172 yes (0.50432277 0.49567723)
##    194) loan_amount< 0.5606189 233 85 yes (0.63519313 0.36480687)
##    388) installment>=0.2022886 43 0 yes (1.00000000 0.00000000) *
##    389) installment< 0.2022886 190 85 yes (0.55263158 0.44736842)
##    778) loan_amount< -0.327277 36 2 yes (0.94444444 0.05555556) *
##    779) loan_amount>=-0.327277 154 71 no (0.46103896 0.53896104)
##   1558) installment>=-0.2936345 100 38 yes (0.62000000 0.38000000)
##   3116) loan_amount< 0.3146527 77 16 yes (0.79220779 0.20779221) *
##   3117) loan_amount>=0.3146527 23 1 no (0.04347826 0.95652174) *
##   1559) installment< -0.2936345 54 9 no (0.16666667 0.83333333) *
##   195) loan_amount>=0.5606189 114 27 no (0.23684211 0.76315789)
##   390) installment>=0.5357504 69 26 no (0.37681159 0.62318841)
##   780) loan_amount< 1.012413 25 3 yes (0.88000000 0.12000000) *
##   781) loan_amount>=1.012413 44 4 no (0.09090909 0.90909091) *
##   391) installment< 0.5357504 45 1 no (0.02222222 0.97777778) *
##   49) loan_amount>=1.59353 38 5 no (0.13157895 0.86842105) *
##  25) installment< -0.8083194 24 0 no (0.00000000 1.00000000) *
##  13) term_five_year< 0.5 814 207 no (0.25429975 0.74570025)
##  26) interest_rate< -0.3269675 13 0 yes (1.00000000 0.00000000) *
##  27) interest_rate>=-0.3269675 801 194 no (0.24219725 0.75780275)
##   54) loan_purpose_medical>=0.5 128 57 no (0.44531250 0.55468750)
##  108) debt_to_income>=-0.03797593 60 20 yes (0.66666667 0.33333333)
##  216) annual_income< 0.1502221 37 8 yes (0.78378378 0.21621622) *
##  217) annual_income>=0.1502221 23 11 no (0.47826087 0.52173913) *
##  109) debt_to_income< -0.03797593 68 17 no (0.25000000 0.75000000) *
##  55) loan_purpose_medical< 0.5 673 137 no (0.20356612 0.79643388)
##  110) loan_purpose_credit_card>=0.5 146 55 no (0.37671233 0.62328767)
##  220) loan_amount< -0.9609695 48 19 yes (0.60416667 0.39583333)
##  440) interest_rate< 0.310133 30 8 yes (0.73333333 0.26666667) *
##  441) interest_rate>=0.310133 18 7 no (0.38888889 0.61111111) *
##  221) loan_amount>=-0.9609695 98 26 no (0.26530612 0.73469388)
##  442) installment>=0.08568682 50 22 no (0.44000000 0.56000000)
##  884) debt_to_income>=0.1172241 30 12 yes (0.60000000 0.40000000) *
##  885) debt_to_income< 0.1172241 20 4 no (0.20000000 0.80000000) *
##  443) installment< 0.08568682 48 4 no (0.08333333 0.91666667) *
##  111) loan_purpose_credit_card< 0.5 527 82 no (0.15559772 0.84440228)
##  222) debt_to_income>=0.9236173 81 24 no (0.29629630 0.70370370) *
##
## ...
## and 10 more lines.
```

```
# Exploring our Trained Model
# Extract the trained model from our workflow fit
tree_fit <- tree_wf_fit %>%
  pull_workflow_fit()
tree_fit
```

```
## parsnip model object
##
```

```


## Fit time: 80ms
## n= 3083
##
## node), split, n, loss, yval, (yprob)
##      * denotes terminal node
##
## 1) root 3083 1148 no (0.37236458 0.62763542)
##    2) interest_rate>=0.691997 702    0 yes (1.00000000 0.00000000) *
##    3) interest_rate< 0.691997 2381  446 no (0.18731625 0.81268375)
##    6) interest_rate>=-0.3606202 1287  446 no (0.34654235 0.65345765)
##    12) term_five_year>=0.5 473   234 yes (0.50528541 0.49471459)
##    24) installment>=-0.8083194 449   210 yes (0.53229399 0.46770601)
##    48) loan_amount< 1.59353 411   177 yes (0.56934307 0.43065693)
##    96) installment>=0.9637937 64     5 yes (0.92187500 0.07812500) *
##    97) installment< 0.9637937 347   172 yes (0.50432277 0.49567723)
##    194) loan_amount< 0.5606189 233   85 yes (0.63519313 0.36480687)
##    388) installment>=0.2022886 43     0 yes (1.00000000 0.00000000) *
##    389) installment< 0.2022886 190   85 yes (0.55263158 0.44736842)
##    778) loan_amount< -0.327277 36     2 yes (0.94444444 0.05555556) *
##    779) loan_amount>=-0.327277 154   71 no (0.46103896 0.53896104)
##    1558) installment>=-0.2936345 100   38 yes (0.62000000 0.38000000)
##    3116) loan_amount< 0.3146527 77    16 yes (0.79220779 0.20779221) *
##    3117) loan_amount>=0.3146527 23     1 no (0.04347826 0.95652174) *
##    1559) installment< -0.2936345 54     9 no (0.16666667 0.83333333) *
##    195) loan_amount>=0.5606189 114   27 no (0.23684211 0.76315789)
##    390) installment>=0.5357504 69    26 no (0.37681159 0.62318841)
##    780) loan_amount< 1.012413 25     3 yes (0.88000000 0.12000000) *
##    781) loan_amount>=1.012413 44     4 no (0.09090909 0.90909091) *
##    391) installment< 0.5357504 45     1 no (0.02222222 0.97777778) *
##    49) loan_amount>=1.59353 38     5 no (0.13157895 0.86842105) *
##    25) installment< -0.8083194 24     0 no (0.00000000 1.00000000) *
##    13) term_five_year< 0.5 814   207 no (0.25429975 0.74570025)
##    26) interest_rate< -0.3269675 13     0 yes (1.00000000 0.00000000) *
##    27) interest_rate>=-0.3269675 801  194 no (0.24219725 0.75780275)
##    54) loan_purpose_medical>=0.5 128    57 no (0.44531250 0.55468750)
##    108) debt_to_income>=-0.03797593 60   20 yes (0.66666667 0.33333333)
##    216) annual_income< 0.1502221 37     8 yes (0.78378378 0.21621622) *
##    217) annual_income>=0.1502221 23    11 no (0.47826087 0.52173913) *
##    109) debt_to_income< -0.03797593 68   17 no (0.25000000 0.75000000) *
##    55) loan_purpose_medical< 0.5 673   137 no (0.20356612 0.79643388)
##    110) loan_purpose_credit_card>=0.5 146   55 no (0.37671233 0.62328767)
##    220) loan_amount< -0.9609695 48    19 yes (0.60416667 0.39583333)
##    440) interest_rate< 0.310133 30     8 yes (0.73333333 0.26666667) *
##    441) interest_rate>=0.310133 18     7 no (0.38888889 0.61111111) *
##    221) loan_amount>=-0.9609695 98    26 no (0.26530612 0.73469388)
##    442) installment>=0.08568682 50    22 no (0.44000000 0.56000000)
##    884) debt_to_income>=0.1172241 30    12 yes (0.60000000 0.40000000) *
##    885) debt_to_income< 0.1172241 20     4 no (0.20000000 0.80000000) *
##    443) installment< 0.08568682 48     4 no (0.08333333 0.91666667) *
##    111) loan_purpose_credit_card< 0.5 527   82 no (0.15559772 0.84440228)
##    222) debt_to_income>=0.9236173 81    24 no (0.29629630 0.70370370) *
##    223) debt_to_income< 0.9236173 446   58 no (0.13004484 0.86995516)
##    446) years_credit_history< -0.3443376 184   34 no (0.18478261 0.81521739)
##    892) current_job_years>=0.004774538 61   18 no (0.29508197 0.70491803)

```

```
##          1784) annual_income< 0.3357285 46    18 no (0.39130435 0.60869565)
##          3568) homeownership_rent>=0.5 28    13 yes (0.53571429 0.46428571) *
##          3569) homeownership_rent< 0.5 18     3 no (0.16666667 0.83333333) *
##          1785) annual_income>=0.3357285 15     0 no (0.00000000 1.00000000) *
##          893) current_job_years< 0.004774538 123    16 no (0.13008130 0.86991870) *
##          447) years_credit_history>=-0.3443376 262    24 no (0.09160305 0.90839695) *
##          7) interest_rate< -0.3606202 1094     0 no (0.00000000 1.00000000) *
```

```
# The variable importance scores from the model
```

```
vip(tree_fit)
```



```
# Predicted Categories
```

```
tree_predictions_categories <- predict(tree_wf_fit,
                                         new_data = loans_test)
tree_predictions_categories
```

```
## # A tibble: 1,027 x 1
##   .pred_class
##   <fct>
## 1 no
## 2 no
## 3 no
## 4 no
## 5 no
## 6 yes
## 7 yes
## 8 yes
```



```

## 9 no
## 10 no
## # ... with 1,017 more rows

# obtain the estimated probabilities for each category of our response variable
tree_predictions_probabilities <- predict(tree_wf_fit,
                                         new_data = loans_test,
                                         type = 'prob')

tree_predictions_probabilities

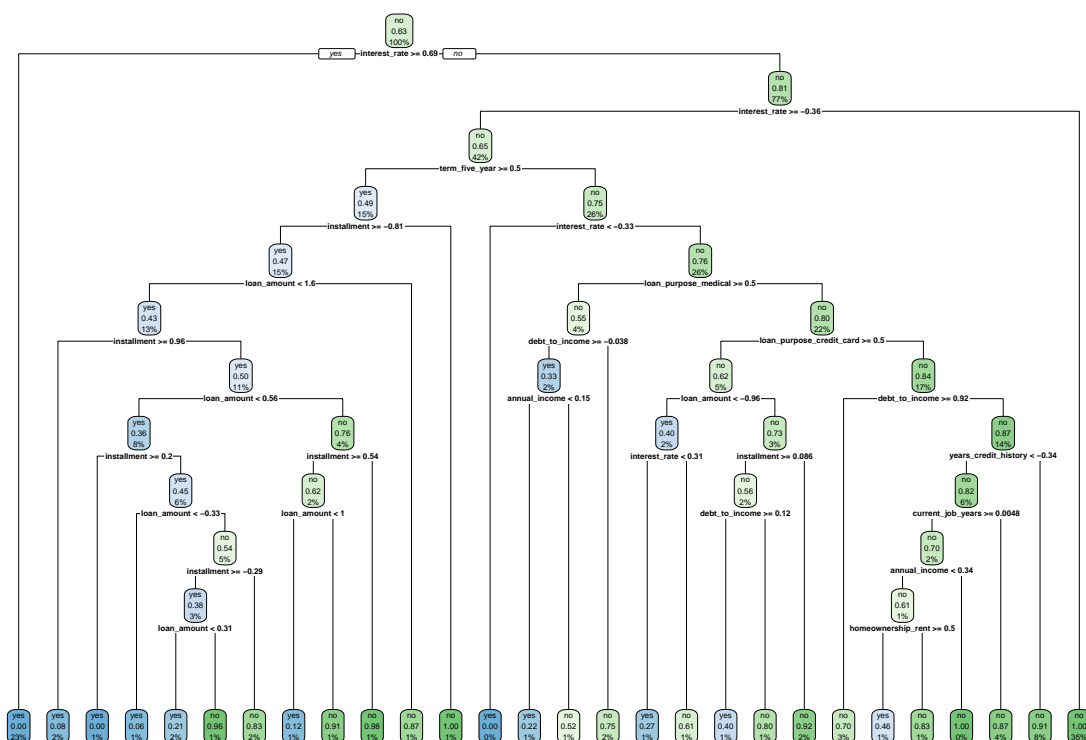
## # A tibble: 1,027 x 2
##   .pred_yes .pred_no
##   <dbl>    <dbl>
## 1      0      1
## 2    0.25    0.75
## 3      0      1
## 4    0.132   0.868
## 5    0.389   0.611
## 6      1      0
## 7      1      0
## 8      1      0
## 9      0      1
## 10     0      1
## # ... with 1,017 more rows

# combine the results from the Predicted Categories and the estimated probabilities with the true resp
tree_test_results <- loans_test %>%
  select(loan_default) %>%
  bind_cols(tree_predictions_categories) %>%
  bind_cols(tree_predictions_probabilities)
tree_test_results

## # A tibble: 1,027 x 4
##   loan_default .pred_class .pred_yes .pred_no
##   <fct>       <fct>       <dbl>    <dbl>
## 1 no         no         0      1
## 2 yes        no         0.25    0.75
## 3 no         no         0      1
## 4 no         no         0.132   0.868
## 5 no         no         0.389   0.611
## 6 yes        yes         1      0
## 7 yes        yes         1      0
## 8 yes        yes         1      0
## 9 no         no         0      1
## 10 no        no         0      1
## # ... with 1,017 more rows

# Decision Tree Plot
rpart.plot(tree_fit$fit,
            roundint = FALSE)

```



#6. Evaluate model performance on the test set by plotting an ROC curve using `autoplot()` and calculating

 # Evaluate the decision tree model

Exploring Performance Metrics

Confusion Matrix

```
conf_mat(tree_test_results,
          truth = loan_default,
          estimate = .pred_class)
```

```
##           Truth
## Prediction yes  no
##           yes 330 28
##           no  52 617
```

F1 Score

```
f_meas(tree_test_results,
        truth = loan_default,
        estimate = .pred_class)
```

```
## # A tibble: 1 x 3
##   .metric .estimator .estimate
##   <chr>   <chr>      <dbl>
## 1 f_meas binary      0.892
```


ROC Curve

```
roc_curve(tree_test_results,
           truth = loan_default,
```

```
estimate = .pred_yes)
```

```
## # A tibble: 25 x 3
##   .threshold specificity sensitivity
##   <dbl>         <dbl>         <dbl>
## 1  -Inf             0             1
## 2    0             0             1
## 3  0.0222         0.592         1
## 4  0.0435         0.606         0.995
## 5  0.0833         0.617         0.995
## 6  0.0909         0.640         0.984
## 7  0.0916         0.656         0.982
## 8  0.130          0.791         0.953
## 9  0.132          0.847         0.937
## 10 0.167          0.862         0.935
## # ... with 15 more rows
```

```
# plot the ROC Curve
roc_curve(tree_test_results,
  truth = loan_default,
  estimate = .pred_yes) %>%
  autoplot()
```



```
# Area Under the ROC Curve
roc_auc(tree_test_results,
  truth = loan_default, .pred_yes)
```

```
## # A tibble: 1 x 3
##   .metric .estimator .estimate
##   <chr>   <chr>       <dbl>
## 1 roc_auc binary      0.973

# Creating Custom Metric Sets
# calculate the accuracy and F1 from my results data frame
my_metrics_tree <- metric_set(accuracy,
                              f_meas)
my_metrics_tree(tree_test_results,
                truth = loan_default,
                estimate = .pred_class)

## # A tibble: 2 x 3
##   .metric .estimator .estimate
##   <chr>   <chr>       <dbl>
## 1 accuracy binary      0.922
## 2 f_meas  binary      0.892

# Automating the Process
tree_last_fit <- final_tree_workflow %>%
  last_fit(loans_split)
tree_last_fit

## # Resampling results
## # Monte Carlo cross-validation (0.75/0.25) with 1 resamples
## # A tibble: 1 x 6
##   splits      id      .metrics    .notes    .predictions    .workflow
##   <list>      <chr>    <list>    <list>    <list>          <list>
## 1 <split [3.1K~ train/test ~ <tibble [2 ~ <tibble [0~ <tibble [1,027 ~ <workflo~

# obtain the metrics on the test set (accuracy and roc_auc by default)
tree_last_fit %>%
  collect_metrics()


## # A tibble: 2 x 3
##   .metric .estimator .estimate
##   <chr>   <chr>       <dbl>
## 1 accuracy binary      0.922
## 2 roc_auc  binary      0.973

# obtain a data frame with test set results
tree_last_fit_results <- tree_last_fit %>%
  collect_predictions()
tree_last_fit_results

## # A tibble: 1,027 x 6
##   id      .pred_yes .pred_no .row .pred_class loan_default
##   <chr>      <dbl>   <dbl> <int> <fct>       <fct>
## 1 train/test split    0       1     3 no         no
## 2 train/test split  0.25    0.75   4 no         yes
## 3 train/test split    0       1     5 no         no
## 4 train/test split  0.132   0.868  10 no         no
## 5 train/test split  0.389   0.611  28 no         no
## 6 train/test split    1       0    33 yes        yes
## 7 train/test split    1       0    34 yes        yes
## 8 train/test split    1       0    48 yes        yes
```

```
## 9 train/test split    0      1      52 no      no
## 10 train/test split   0      1      69 no      no
## # ... with 1,017 more rows

# plot the ROC curve to visualize test set performance of the tuned decision tree
tree_last_fit_results %>%
  roc_curve(truth = loan_default,
            estimate = .pred_yes) %>%
  autoplot()
```



Summary of Results [25 Points]

Write a summary of your overall findings and recommendations to the executives at the bank. Think of this section as your closing remarks of a presentation, where you summarize your key findings, model performance, and make recommendations to improve loan processes at the bank.

Your summary should include:

- Key findings from your data analysis. What were the things that stuck out for you in this section and why are they important?
- Your “best” classification model and an analysis of its performance.
 - In this section you should talk about the expected error of your model on future data.
 - You should discuss at least one performance metric, such as an F1 or ROC AUC for your model. However, you must explain the results in an intuitive, non-technical manner. Your audience in this case are executives at a bank with limited knowledge of machine learning.
- Your recommendations to the bank – how can financial losses from default be reduced?

Summary

The main thing that can be concluded by the data analysis part is that the number of the total number of the non default customer is higher than the loan default customer. Therefore, problem might occur if there is a special rare case for the loan default customer appears. Other than that, there is some predictors which would suppose to be better if provided by a set fix of value or factor instead of being the numeric type. For an instance, the interest rate which might be better to be replaced by type of loan product so that the difference between the interest rate and the loan default can be easily found since each loan product should usually has a fix interest rate. Also, information about loan collectibility should be provided as well in the dataset because the relationship with the loan default is usually strong.

The classifications models, which are used to predict the loan default customer, are logistic regression and decision tree. Those two models are tested through the test set which were obtained from the first split at beginning. The number of row for the test set is 1027. Therefore, the two models are tested by using the same 1027 rows dataset.

Firstly, we will explore the performance through the confusion metric. Respectively, the confusion metric of the first and the second model are such the follows:

- logistic regression: Truth Prediction yes no yes 354 16 no 28 629
- decision tree Truth Prediction yes no yes 330 28 no 52 617

According to the first confusion metric result, we can see that there are 983 rows out of 1027 have the right prediction. Specifically, in that 983 rows, there 354 rows of the customers loan default have the same value with the predicted value and there are 629 of the non default customers have the same value with the predicted value. However, there are 44 rows which are missed predicted since the actual values are different to the predicted value. Therefore, according to such confusion metric, we can conclude that the accuracy of this model is 95%.

However, according to the second confusion metric, we can calculate that the accuracy of the model is 92%. Therefore, the accuracy of the second model, which is the decision tree model, based on the confusion metric, is not as good as the first model. Yet, the difference of the two wrong predicted values of the second model are quite significant, which are 52 and 28. Therefore, the F1 score should be calculated in order to support the performance measurement of the second model.

According to the both accuracy, the expected error from the test set respectively of those two models are 0.043 and 0.078. In other words, the tendency of the logistic regression model to yield the error result is 4.3% and the decision tree is 7.8%. Therefore, the logistic regression still has the smaller chance to gain error results while predicting whether the customer tend to be default or non default in loan payment.

Secondly, we explore the F1 score that serves as the performance metric which balances the missed predicted values including the customers which are wrongly predicted as the loan default or the non default status. Respectively, the F1 score of the first model and second model while tested by using the test set is such the follows:

- logistic regression: f₁ meas binary 0.9414894
- decision tree: f₁ meas binary 0.8918919

According to number of missed predicted value in the confusion metric of the first model above, the difference between false positive, which is the wrong prediction for the default customer, and the false negative, which is the total number of the wrong prediction for the non default customer, are not significantly different, using the accuracy 95% accuracy as the performance measurement is still okay. Even though, according to the F1 score, the performance value is still very good, which is 0.94, since the value is close to 1. Obviously, compared to the F1 score of the second model, the first model still has the better performance since the F1 score of the second model is 0.89 which is lower than the F1 score of the logistic regression model.

Thirdly, we explore the ROC curve to measure the model performance. According to the ROC Curve of the logistic regression model, the plot between the sensitivity, which is the true positive fraction (the true “yes” of

loan default), and the specificity, which is the false positive fraction, is close to the 0,1 point. The first model is even better than the second model since based on the ROC curve, the first model seems to have the closer curve to the 0,1 point, which is on top left corner.

The last one, we explore the Area Under the ROC Curve. Respectively, the AU ROC of the Logistic model and the decision tree are such the follows:

- logistic regression: roc_auc binary 0.9894598
- decision tree: roc_auc binary 0.9729595

The AU ROC, which falls between 0.9 to 1, shows a good model performance. Since the AU ROC for the logistic regression model is 0.9894598, which is very close to 1, then the model is classified as a very good model. The same as the decision tree model, the both models can be considered as the good model. However, in overall, the logistic regression model is the better model to the classification for the loan default customer based on the given dataset.

According to all the result, I suggest the bank to pay attention more to the customers as the following criterias:

1. The customers who use the loan for credit card and medical purpose
2. The customers within 5 years term of loan
3. The customers who have worked for 8 years in the current job
4. The customers which have the loan amount above 10000
5. The customers which have installment above 490.00
6. The customers who live in rent houses

The above criterias shows that the customers have tendency to get the loan default status. Therefore, those kind of customers can be treated by restructuring the loan account whether by extending the payment term and reducing the interest rate, or reverting the method of loan payment by paying principle loan amount and followed by interest payment, or vice versa.